

CNNFruit

May 18, 2022

```
[53]: import tensorflow as tf
      from tensorflow import keras
      from keras.models import Sequential
      from keras.layers.convolutional import Conv2D, MaxPooling2D
      from keras.layers import Flatten, Dense, Dropout, Activation
      from google.colab import drive
```

```
[54]: drive.mount('/content/drive',force_remount=True)
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      train_datagen = ImageDataGenerator(rescale=1./255,
                                         shear_range=0.2,
                                         zoom_range=0.2,
                                         horizontal_flip=True)
      train=train_datagen.flow_from_directory('/content/drive/MyDrive/train',
                                              target_size=(256,256),
                                              batch_size=32,
                                              class_mode='categorical')
      test=train_datagen.flow_from_directory('/content/drive/MyDrive/test1',
                                              target_size=(256,256),
                                              batch_size=32,
                                              class_mode='categorical')
```

```
Mounted at /content/drive
Found 48 images belonging to 10 classes.
Found 20 images belonging to 10 classes.
```

```
[55]: drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).
```

```
[56]: model=Sequential()
      model.
      ↪add(Conv2D(128,(3,3),activation='relu',kernel_initializer='he_uniform',padding='same',input
      model.add(MaxPooling2D(pool_size=(2,2)))
      model.
      ↪add(Conv2D(64,(3,3),activation='relu',kernel_initializer='he_uniform',padding='same'))
      model.add(MaxPooling2D((2,2)))
```

```

model.add(Conv2D(32,(3,3),activation='relu',kernel_initializer='he_uniform',padding='same'))
model.add(MaxPooling2D((2,2)))
model.add(Flatten())
model.add(Dense(128,activation='relu',kernel_initializer = 'he_uniform'))
#model.add(Dropout(0.2))
model.add(Dense(10,activation='Softmax'))
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.callbacks import EarlyStopping
#opt = SGD(lr = 0.01, momentum = 0.9)
model.compile(optimizer = 'adam', loss = 'categorical_crossentropy',metrics =_
    →['accuracy'])
callbacks=[EarlyStopping(monitor='val_loss',patience=100)]
history=model.fit(train,
                    steps_per_epoch=len(train),
                    batch_size = 64,
                    epochs=50,
                    validation_data=test,
                    validation_steps=len(test),
                    callbacks=callbacks,
                    verbose = 1)

```

Epoch 1/50

2/2 [=====] - 7s 5s/step - loss: 26.9320 - accuracy:
0.0833 - val_loss: 22.7313 - val_accuracy: 0.1500

Epoch 2/50

2/2 [=====] - 5s 2s/step - loss: 19.0183 - accuracy:
0.1250 - val_loss: 6.1188 - val_accuracy: 0.1000

Epoch 3/50

2/2 [=====] - 3s 2s/step - loss: 3.8251 - accuracy:
0.1042 - val_loss: 2.8134 - val_accuracy: 0.1000

Epoch 4/50

2/2 [=====] - 2s 985ms/step - loss: 2.5122 - accuracy:
0.1458 - val_loss: 2.2533 - val_accuracy: 0.2000

Epoch 5/50

2/2 [=====] - 2s 1s/step - loss: 2.2528 - accuracy:
0.1042 - val_loss: 2.2392 - val_accuracy: 0.1500

Epoch 6/50

2/2 [=====] - 2s 2s/step - loss: 2.2668 - accuracy:
0.1667 - val_loss: 2.2018 - val_accuracy: 0.2000

Epoch 7/50

2/2 [=====] - 4s 3s/step - loss: 2.0472 - accuracy:
0.3542 - val_loss: 2.0598 - val_accuracy: 0.2000

Epoch 8/50

2/2 [=====] - 2s 2s/step - loss: 2.3673 - accuracy:
0.2500 - val_loss: 2.2420 - val_accuracy: 0.1000

Epoch 9/50

2/2 [=====] - 2s 2s/step - loss: 2.1531 - accuracy:
0.1875 - val_loss: 2.1130 - val_accuracy: 0.2500
Epoch 10/50
2/2 [=====] - 2s 1s/step - loss: 1.8890 - accuracy:
0.3750 - val_loss: 1.9804 - val_accuracy: 0.2500
Epoch 11/50
2/2 [=====] - 2s 2s/step - loss: 1.7227 - accuracy:
0.4167 - val_loss: 1.8396 - val_accuracy: 0.2500
Epoch 12/50
2/2 [=====] - 2s 2s/step - loss: 1.4337 - accuracy:
0.4792 - val_loss: 1.6758 - val_accuracy: 0.3000
Epoch 13/50
2/2 [=====] - 2s 996ms/step - loss: 1.2073 - accuracy:
0.6250 - val_loss: 1.8741 - val_accuracy: 0.2000
Epoch 14/50
2/2 [=====] - 2s 1s/step - loss: 1.0129 - accuracy:
0.6250 - val_loss: 1.9103 - val_accuracy: 0.3000
Epoch 15/50
2/2 [=====] - 2s 2s/step - loss: 0.7869 - accuracy:
0.7917 - val_loss: 1.7886 - val_accuracy: 0.3500
Epoch 16/50
2/2 [=====] - 2s 2s/step - loss: 0.6814 - accuracy:
0.8125 - val_loss: 1.6292 - val_accuracy: 0.3000
Epoch 17/50
2/2 [=====] - 2s 1s/step - loss: 0.5875 - accuracy:
0.8333 - val_loss: 1.9560 - val_accuracy: 0.3500
Epoch 18/50
2/2 [=====] - 2s 2s/step - loss: 0.4534 - accuracy:
0.8750 - val_loss: 1.6596 - val_accuracy: 0.3000
Epoch 19/50
2/2 [=====] - 2s 2s/step - loss: 0.4009 - accuracy:
0.9167 - val_loss: 1.9364 - val_accuracy: 0.3500
Epoch 20/50
2/2 [=====] - 2s 1s/step - loss: 0.2458 - accuracy:
0.9583 - val_loss: 1.7112 - val_accuracy: 0.5000
Epoch 21/50
2/2 [=====] - 2s 1s/step - loss: 0.3895 - accuracy:
0.8542 - val_loss: 1.8460 - val_accuracy: 0.5000
Epoch 22/50
2/2 [=====] - 2s 2s/step - loss: 0.1189 - accuracy:
1.0000 - val_loss: 1.7692 - val_accuracy: 0.4500
Epoch 23/50
2/2 [=====] - 2s 1s/step - loss: 0.1110 - accuracy:
1.0000 - val_loss: 1.6497 - val_accuracy: 0.3500
Epoch 24/50
2/2 [=====] - 2s 2s/step - loss: 0.0992 - accuracy:
0.9792 - val_loss: 2.4550 - val_accuracy: 0.4500
Epoch 25/50

2/2 [=====] - 2s 2s/step - loss: 0.1386 - accuracy:
 0.9583 - val_loss: 2.0278 - val_accuracy: 0.5000
 Epoch 26/50
 2/2 [=====] - 2s 2s/step - loss: 0.1242 - accuracy:
 0.9583 - val_loss: 1.8718 - val_accuracy: 0.3500
 Epoch 27/50
 2/2 [=====] - 2s 2s/step - loss: 0.1125 - accuracy:
 0.9792 - val_loss: 2.5232 - val_accuracy: 0.3500
 Epoch 28/50
 2/2 [=====] - 2s 1s/step - loss: 0.2167 - accuracy:
 0.9167 - val_loss: 2.2803 - val_accuracy: 0.4000
 Epoch 29/50
 2/2 [=====] - 3s 1s/step - loss: 0.1482 - accuracy:
 0.9583 - val_loss: 2.5548 - val_accuracy: 0.4500
 Epoch 30/50
 2/2 [=====] - 2s 2s/step - loss: 0.1397 - accuracy:
 0.9375 - val_loss: 2.3721 - val_accuracy: 0.3500
 Epoch 31/50
 2/2 [=====] - 2s 2s/step - loss: 0.2261 - accuracy:
 0.9583 - val_loss: 2.2710 - val_accuracy: 0.3000
 Epoch 32/50
 2/2 [=====] - 2s 2s/step - loss: 0.2396 - accuracy:
 0.9375 - val_loss: 2.2715 - val_accuracy: 0.5000
 Epoch 33/50
 2/2 [=====] - 2s 1s/step - loss: 0.0762 - accuracy:
 0.9583 - val_loss: 2.7368 - val_accuracy: 0.4500
 Epoch 34/50
 2/2 [=====] - 2s 2s/step - loss: 0.1704 - accuracy:
 0.9583 - val_loss: 3.0861 - val_accuracy: 0.3500
 Epoch 35/50
 2/2 [=====] - 2s 2s/step - loss: 0.2951 - accuracy:
 0.8750 - val_loss: 2.9272 - val_accuracy: 0.4000
 Epoch 36/50
 2/2 [=====] - 2s 2s/step - loss: 0.2649 - accuracy:
 0.9167 - val_loss: 2.4136 - val_accuracy: 0.4000
 Epoch 37/50
 2/2 [=====] - 2s 2s/step - loss: 0.7057 - accuracy:
 0.7708 - val_loss: 1.6669 - val_accuracy: 0.4000
 Epoch 38/50
 2/2 [=====] - 2s 1s/step - loss: 0.2149 - accuracy:
 0.9375 - val_loss: 2.5152 - val_accuracy: 0.4000
 Epoch 39/50
 2/2 [=====] - 2s 2s/step - loss: 0.2266 - accuracy:
 0.9167 - val_loss: 3.4246 - val_accuracy: 0.3500
 Epoch 40/50
 2/2 [=====] - 2s 2s/step - loss: 0.4013 - accuracy:
 0.8750 - val_loss: 2.2826 - val_accuracy: 0.4500
 Epoch 41/50

```

2/2 [=====] - 2s 2s/step - loss: 0.3389 - accuracy:
0.9167 - val_loss: 1.8615 - val_accuracy: 0.4500
Epoch 42/50
2/2 [=====] - 2s 2s/step - loss: 0.0846 - accuracy:
0.9792 - val_loss: 1.9804 - val_accuracy: 0.5500
Epoch 43/50
2/2 [=====] - 2s 1s/step - loss: 0.0419 - accuracy:
1.0000 - val_loss: 2.7068 - val_accuracy: 0.4000
Epoch 44/50
2/2 [=====] - 2s 2s/step - loss: 0.1346 - accuracy:
0.9792 - val_loss: 2.6812 - val_accuracy: 0.4500
Epoch 45/50
2/2 [=====] - 2s 2s/step - loss: 0.0623 - accuracy:
0.9792 - val_loss: 3.1316 - val_accuracy: 0.4000
Epoch 46/50
2/2 [=====] - 2s 2s/step - loss: 0.1741 - accuracy:
0.9792 - val_loss: 3.1748 - val_accuracy: 0.3500
Epoch 47/50
2/2 [=====] - 2s 1s/step - loss: 0.0872 - accuracy:
0.9792 - val_loss: 3.1480 - val_accuracy: 0.3000
Epoch 48/50
2/2 [=====] - 2s 2s/step - loss: 0.1012 - accuracy:
0.9792 - val_loss: 2.7172 - val_accuracy: 0.3500
Epoch 49/50
2/2 [=====] - 2s 1s/step - loss: 0.0459 - accuracy:
0.9792 - val_loss: 2.5316 - val_accuracy: 0.4500
Epoch 50/50
2/2 [=====] - 2s 2s/step - loss: 0.0075 - accuracy:
1.0000 - val_loss: 2.7316 - val_accuracy: 0.4000

```

```

[57]: #đánh giá chất lượng của mô hình và vẽ lại
score = model.evaluate(test,verbose=0)
print('Sai số kiểm tra là: ',score[0])
print('Độ chính xác kiểm tra là: ',score[1])

```

```

Sai số kiểm tra là: 2.5663399696350098
Độ chính xác kiểm tra là: 0.44999998807907104

```

```

[58]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epochs')
plt.legend(['train','validation'],loc='upper_left')

```

```


/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6:
MatplotlibDeprecationWarning: Unrecognized location 'upper_left'. Falling back
on 'best'; valid locations are

```

```
best
upper right
upper left
lower left
lower right
right
center left
center right
lower center
upper center
center
```

This will raise an exception in 3.3.

[58]: <matplotlib.legend.Legend at 0x7f119f783190>



```
[59]: model.save('model_fruit.h5')
```

```
[60]: from tensorflow.keras.models import load_model
model=load_model('model_fruit.h5')
```

```
[61]: from tensorflow.keras.utils import load_img
from tensorflow.keras.utils import img_to_array
import matplotlib.pyplot as plt
```

```
import pandas as pd
import numpy as np
import math
```


```
[79]: from keras.preprocessing.image import load_img
      from keras.preprocessing.image import img_to_array
      from google.colab import files
      uploadfile=files.upload()
```

<IPython.core.display.HTML object>

Saving chuoi.jpg to chuoi (1).jpg

```
[64]: fruit = ['bưởi', 'cam', 'chuối', 'kiwi', 'mận', 'mít', 'nhãn', 'táo', 'thanh_
      ↳long', 'xoài']
```

```
[80]: img = load_img("chuoi.jpg", target_size = (256, 256))
      plt.imshow(img)
      img=img_to_array(img)
      img=img.reshape(1,256,256,3)
      img=img.astype('float32')
      img=img/255
```



```
[81]: np.argmax(model.predict(img), axis=1)
```

```
[81]: array([2])
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
!wget -nc https://raw.githubusercontent.com/brpy/colab-pdf/master/colab_pdf.py
from colab_pdf import colab_pdf
colab_pdf('CNNFruit.ipynb')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

--2022-05-18 07:43:52-- https://raw.githubusercontent.com/brpy/colab-pdf/master/colab_pdf.py

Resolving raw.githubusercontent.com (raw.githubusercontent.com)...

185.199.108.133, 185.199.109.133, 185.199.110.133, ...

Connecting to raw.githubusercontent.com

(raw.githubusercontent.com)|185.199.108.133|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 1864 (1.8K) [text/plain]

Saving to: 'colab_pdf.py'

colab_pdf.py 100%[=====>] 1.82K --.-KB/s in 0s

2022-05-18 07:43:53 (27.8 MB/s) - 'colab_pdf.py' saved [1864/1864]

WARNING: apt does not have a stable CLI interface. Use with caution in scripts.

WARNING: apt does not have a stable CLI interface. Use with caution in scripts.

Extracting templates from packages: 100%

```
[ ]:
```